

# COMPSCI 590N

## Lecture 7: Complexity and Numerical Linear Algebra 1

Roy J. Adams

College of Information and Computer Sciences  
University of Massachusetts Amherst

# Outline

## 1 Assignment 2


## 2 NumPy Views

## 3 Computational Complexity

## 4 Numerical Linear Algebra


# Problem 1

```
crsl = Course(589)
```



# Problem 1

```
crs1 = Course(589)
std1 = Student(123)
```



# Problem 1

```
crs1 = Course(589)
sdt1 = Student(123)
crs1.add_student(sdt1)
```

crs1


id = 589  
roster = {}

sdt1

id = 123  
crs = set([])

# Problem 1

```
crs1 = Course(589)
std1 = Student(123)
crs1.add_student(std1)
```





---

```
self.roster[student.id] = student
```

# Problem 1

```
crs1 = Course(589)
std1 = Student(123)
crs1.add_student(std1)
```





---

```
self.roster[student.id] = student
student.roster.add(self.id)
```

# Problem 1


```
crs1 = Course(589)
sdt1 = Student(123)
crs1.add_student(sdt1)
crs1.drop_student(123)
```





# Problem 1

```
crs1 = Course(589)
sdt1 = Student(123)
crs1.add_student(sdt1)
crs1.drop_student(123)
```





---

```
student = self.roster[student_id]
```

# Problem 1

```
crs1 = Course(589)
sdt1 = Student(123)
crs1.add_student(sdt1)
crs1.drop_student(123)
```





---

```
student = self.roster[student_id]
student.crs.remove(self.id)
```

# Problem 1

```
crs1 = Course(589)
std1 = Student(123)
crs1.add_student(std1)
crs1.drop_student(123)
```




---

```
student = self.roster[student_id]
student.crs.remove(self.id)
del self.roster[student_id]
```

# Outline

- 1 Assignment 2
- 2 NumPy Views
- 3 Computational Complexity
- 4 Numerical Linear Algebra

# Views

A `view` of an array is another `ndarray` object that shares the same underlying memory locations.

# Views

A view of an array is another ndarray object that shares the same underlying memory locations.

```
>>> A = np.ones((4,5))
>>> B = A.view()
>>> B is A
False
>>> B.base is A
True
```

# Views

A view of an array is another ndarray object that shares the same underlying memory locations.

```
>>> A = np.ones((4,5))
>>> B = A.view()
>>> B is A
False
>>> B.base is A
True
```

- A and B are not the same object, hence `A is B` returns `False`.

# Views

A view of an array is another ndarray object that shares the same underlying memory locations.

```
>>> A = np.ones((4,5))
>>> B = A.view()
>>> B is A
False
>>> B.base is A
True
```

- A and B are not the same object, hence `A is B` returns `False`.
- The attribute `base` points to the array from which a view was created, hence `B.base is A` returns `True`.



# Views

Many functions will try to create views if they can, **but will silently create copies if they can't**. Examples of such operations/functions are:

# Views

Many functions will try to create views if they can, **but will silently create copies if they can't**. Examples of such operations/functions are:

- Slicing

# Views

Many functions will try to create views if they can, **but will silently create copies if they can't**. Examples of such operations/functions are:

- Slicing
- transpose

# Views

Many functions will try to create views if they can, **but will silently create copies if they can't**. Examples of such operations/functions are:

- Slicing
- transpose
- reshape

# Views

```
>>> A = np.ones((4,5))
>>> B = A.transpose()
>>> B.base is A
True
>>> C = B.reshape((5,4)) # Creates another view
>>> C.base is A
True
>>> D = B.reshape((2,10)) # Creates a copy!
>>> D.base is A
False
```

# Views

- There are rules for when a view will be made, however, they are not well documented and rely on an understanding of the underlying storage.

# Views

- There are rules for when a view will be made, however, they are not well documented and rely on an understanding of the underlying storage.
- The moral of the story is: be careful when modifying the contents of an array directly.

# Views

- There are rules for when a view will be made, however, they are not well documented and rely on an understanding of the underlying storage.
- The moral of the story is: be careful when modifying the contents of an array directly.
- When in doubt, create a copy!



# Outline

- 1 Assignment 2
- 2 NumPy Views
- 3 Computational Complexity**
- 4 Numerical Linear Algebra

# Computational Complexity

- In our discussion of numerical algorithms so far, we have obliquely discussed the speed of our algorithms.

# Computational Complexity

- In our discussion of numerical algorithms so far, we have obliquely discussed the speed of our algorithms.
- We need a formal language for discussing the speed of an algorithm.

# Computational Complexity

- In our discussion of numerical algorithms so far, we have obliquely discussed the speed of our algorithms.
- We need a formal language for discussing the speed of an algorithm.
- The mathematical language used by computer scientists and mathematicians is called *complexity*.

# Computational Complexity

High level idea: The complexity of an algorithm is the number of *basic operations* used by the algorithm as a function of the input size.

# Computational Complexity

High level idea: The complexity of an algorithm is the number of *basic operations* used by the algorithm as a function of the input size.

## Basic operations

A basic operation is any computation that has a constant runtime as a function of the input size. Think of basic operations as the building blocks of algorithms.

# Computational Complexity

High level idea: The complexity of an algorithm is the number of *basic operations* used by the algorithm as a function of the input size.

## Basic operations

A basic operation is any computation that has a constant runtime as a function of the input size. Think of basic operations as the building blocks of algorithms.

Examples include: arithmetic operations, basic functions, logical comparisons, basic indexing, and assignment.

# Array Scanning

Let's consider an example: Finding the maximum value in a list of numbers.



# Array Scanning

Let's consider an example: Finding the maximum value in a list of numbers.

```
def max(A):  
    m = -inf  
    for a in A:  
        if a > m:  
            m = a  
    return m
```

# Array Scanning

Let's consider an example: Finding the maximum value in a list of numbers.

```
def max(A):  
    m = -inf  
    for a in A:  
        if a > m:  
            m = a  
    return m
```

- What are the basic operations used in this algorithm?

# Array Scanning

Let's consider an example: Finding the maximum value in a list of numbers.

```
def max(A):  
    m = -inf  
    for a in A:  
        if a > m:  
            m = a  
    return m
```

- What are the basic operations used in this algorithm?
- If  $A$  has length  $n$ , how many basic operations will be performed by `max`?

# Array Scanning

Let's consider an example: Finding the maximum value in a list of numbers.

```
def max(A):  
    m = -inf  
    for a in A:  
        if a > m:  
            m = a  
    return m
```

- What are the basic operations used in this algorithm?
- If  $A$  has length  $n$ , how many basic operations will be performed by `max`?
- Aside from  $n$ , does the run time of `max` depend on the contents of  $A$ ?

# Worst case analysis

In the previous, example, `max` didn't depend on the contents of `A` beyond its size. This is not always the case.

# Worst case analysis

In the previous, example, `max` didn't depend on the contents of `A` beyond its size. This is not always the case.

```
def find_first_zero(A):  
    i = 0  
    for a in A:  
        if a == 0:  
            return i  
        else:  
            i += 1  
    return None
```

# Worst case analysis

In the previous, example, `max` didn't depend on the contents of `A` beyond its size. This is not always the case.

```
def find_first_zero(A):  
    i = 0  
    for a in A:  
        if a == 0:  
            return i  
        else:  
            i += 1  
    return None
```

- What are the basic operations used in this algorithm?

# Worst case analysis

In the previous, example, `max` didn't depend on the contents of `A` beyond its size. This is not always the case.

```
def find_first_zero(A):  
    i = 0  
    for a in A:  
        if a == 0:  
            return i  
        else:  
            i += 1  
    return None
```

- What are the basic operations used in this algorithm?
- If `A` has length  $n$ , how many basic operations will be performed by `find_first_zero`?



# Worst case analysis

In the previous, example, `max` didn't depend on the contents of `A` beyond its size. This is not always the case.

```
def find_first_zero(A):  
    i = 0  
    for a in A:  
        if a == 0:  
            return i  
        else:  
            i += 1  
    return None
```

- What are the basic operations used in this algorithm?
- If `A` has length  $n$ , how many basic operations will be performed by `find_first_zero`?
- It is standard practice to assume the worst possible contents of `A`.

This is called **worst-case analysis**.

# Binary Search

Problem: Given a **sorted** list of unique numbers,  $A$ , and a number of interest,  $x$ , find the position of  $x$  in  $A$ . Return `None` if  $x$  is not in  $A$ .

# Binary Search

Problem: Given a **sorted** list of unique numbers,  $A$ , and a number of interest,  $x$ , find the position of  $x$  in  $A$ . Return `None` if  $x$  is not in  $A$ .

```
def linear_search(A, x):  
    i = 0  
    for a in A:  
        if a == x:  
            return i  
        else:  
            i += 1  
    return None
```

# Binary Search

Problem: Given a **sorted** list of unique numbers,  $A$ , and a number of interest,  $x$ , find the position of  $x$  in  $A$ . Return `None` if  $x$  is not in  $A$ .

```
def linear_search(A, x):  
    i = 0  
    for a in A:  
        if a == x:  
            return i  
        else:  
            i += 1  
    return None
```

- What is the worst case?

# Binary Search

Problem: Given a **sorted** list of unique numbers,  $A$ , and a number of interest,  $x$ , find the position of  $x$  in  $A$ . Return `None` if  $x$  is not in  $A$ .

```
def linear_search(A, x):  
    i = 0  
    for a in A:  
        if a == x:  
            return i  
        else:  
            i += 1  
    return None
```

- What is the worst case?
- If  $A$  has length  $n$ , what is the worst case number of operations performed by `linear_search`?

# Binary Search

Problem: Given a **sorted** list of unique numbers,  $A$ , and a number of interest,  $x$ , find the position of  $x$  in  $A$ . Return `None` if  $x$  is not in  $A$ .

```
def linear_search(A, x):  
    i = 0  
    for a in A:  
        if a == x:  
            return i  
        else:  
            i += 1  
    return None
```

- What is the worst case?
- If  $A$  has length  $n$ , what is the worst case number of operations performed by `linear_search`?
- What information is not being used by `linear_search`?

# Binary Search

Instead, we can use a more sophisticated algorithm called **binary search**.

# Binary Search

Instead, we can use a more sophisticated algorithm called **binary search**.

## Binary Search: Basic Idea

- Look at the item in the middle of the list, call this item  $m$ .



# Binary Search

Instead, we can use a more sophisticated algorithm called **binary search**.

## Binary Search: Basic Idea

- Look at the item in the middle of the list, call this item  $m$ .
- We know that everything before  $m$  is less than  $m$  and everything after is greater than  $m$ .

# Binary Search

Instead, we can use a more sophisticated algorithm called **binary search**.

## Binary Search: Basic Idea

- Look at the item in the middle of the list, call this item  $m$ .
- We know that everything before  $m$  is less than  $m$  and everything after is greater than  $m$ .
- Therefore, if  $x > m$ , we can eliminate the first half of the list and if  $x < m$  we can eliminate the second half.

# Binary Search

Instead, we can use a more sophisticated algorithm called **binary search**.

## Binary Search: Basic Idea

- Look at the item in the middle of the list, call this item  $m$ .
- We know that everything before  $m$  is less than  $m$  and everything after is greater than  $m$ .
- Therefore, if  $x > m$ , we can eliminate the first half of the list and if  $x < m$  we can eliminate the second half.
- Apply this idea recursively.

# Binary Search

# Animation

# Binary Search

The following code implements binary search in Python.

# Binary Search

The following code implements binary search in Python.

```
def binary_search(A, x):  
    cur_min_idx = 0  
    cur_max_idx = len(A) - 1  
    while cur_min_idx <= cur_max_idx:  
        cur_split = cur_min_idx + np.floor((cur_max_idx  
            - cur_min_idx)/2)  
        if A[cur_split] > x:  
            cur_min_idx = cur_split + 1  
        elif A[cur_split] < x:  
            cur_max_idx = cur_split - 1  
        else:  
            return cur_split  
    return None
```

# Binary Search

- What is the worst case?

# Binary Search

- What is the worst case?
- If  $A$  has length  $n$ , what is the worst case number of operations performed by `binary_search`?



# Binary Search

- What is the worst case?
- If  $A$  has length  $n$ , what is the worst case number of operations performed by `binary_search`?
  - In each iteration, we are dividing the list in half so `binary_search` will take  $\log(n)$  iterations.

# Big O Notation

Computer scientists and mathematicians have a formal notation for discussing the run time of an algorithm as a function of the algorithm's input size.

# Big O Notation

Computer scientists and mathematicians have a formal notation for discussing the run time of an algorithm as a function of the algorithm's input size.


## Big O Notation

Let  $f$  and  $g$  be functions, then  $f(x) = \mathcal{O}(g(x))$  if and only if there exists a constant  $M$  and  $x_0$  such that

$$|f(x)| \leq M|g(x)| \text{ for all } x \geq x_0$$

Intuitively, this means that for large numbers  $g(x)$  is always greater than  $f(x)$ , up to a constant.

# Big O Notation



# Takeaways

- An algorithm's **complexity** is a measure of its run-time as a function of the input size, ignoring constant factors.

# Takeaways

- An algorithm's **complexity** is a measure of its run-time as a function of the input size, ignoring constant factors.
- It is typical to consider the **worst-case** input of a fixed size.

# Takeaways

- An algorithm's **complexity** is a measure of its run-time as a function of the input size, ignoring constant factors.
- It is typical to consider the **worst-case** input of a fixed size.
- Big-O notation is used to formalize this concept so that algorithms can be rigorously compared.

# Outline

- 1 Assignment 2
- 2 NumPy Views
- 3 Computational Complexity
- 4 Numerical Linear Algebra**



# Intro

- Numerical Linear Algebra is the study of efficient algorithms for performing linear algebra computations.

# Intro

- Numerical Linear Algebra is the study of efficient algorithms for performing linear algebra computations.
- Linear algebra computations are at the core **many** algorithms.  
For example:

# Intro

- Numerical Linear Algebra is the study of efficient algorithms for performing linear algebra computations.
- Linear algebra computations are at the core **many** algorithms. For example:
  - Linear regression uses matrix multiplication and inversion.

# Intro

- Numerical Linear Algebra is the study of efficient algorithms for performing linear algebra computations.
- Linear algebra computations are at the core **many** algorithms. For example:
  - Linear regression uses matrix multiplication and inversion.
  - PageRank, the algorithm that spawned Google, was fundamentally calculating eigenvalues.

# Intro

- Numerical Linear Algebra is the study of efficient algorithms for performing linear algebra computations.
- Linear algebra computations are at the core **many** algorithms. For example:
  - Linear regression uses matrix multiplication and inversion.
  - PageRank, the algorithm that spawned Google, was fundamentally calculating eigenvalues.
  - Many of the recommendation systems that power sites like Netflix, Amazon, Google Ads, etc. are based on matrix decomposition.

# Element-wise Operations

- The most basic operations we have discussed are element-wise operations.

# Element-wise Operations

- The most basic operations we have discussed are element-wise operations.
- Example: What is the complexity of multiplying an  $n \times n$  matrix by a scalar?

# Element-wise Operations

- The most basic operations we have discussed are element-wise operations.
- Example: What is the complexity of multiplying an  $n \times n$  matrix by a scalar?



# Element-wise Operations

- The most basic operations we have discussed are element-wise operations.
- Example: What is the complexity of multiplying an  $n \times n$  matrix by a scalar?

```
Input: A, alpha
Output: B
for i in 1,...,n:
    for j in 1,...,n:
        B[i,j] = alpha*A[i,j]
```

# Element-wise Operations

- The most basic operations we have discussed are element-wise operations.
- Example: What is the complexity of multiplying an  $n \times n$  matrix by a scalar?

```
Input: A, alpha
Output: B
for i in 1,...,n:
    for j in 1,...,n:
        B[i,j] = alpha*A[i,j]
```

Scalar/matrix operations have complexity  $\mathcal{O}(n^2)$  for an  $n \times n$  matrix.

# Element-wise Operations

- What about matrix addition?

# Element-wise Operations

- What about matrix addition?

# Element-wise Operations

## ■ What about matrix addition?

```
Input: A, B
Output: C
for i in 1,...,n:
    for j in 1,...,n:
        C[i,j] = A[i,j] + B[i,j]
```

# Element-wise Operations

## ■ What about matrix addition?

```
Input: A, B
Output: C
for i in 1,...,n:
    for j in 1,...,n:
        C[i,j] = A[i,j] + B[i,j]
```

Element-wise matrix/matrix operations have complexity  $\mathcal{O}(n^2)$  for two  $n \times n$  matrices.

# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.

# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.
- Take for example the inner product of two length  $n$  vectors,  $x$  and  $y$ .



# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.
- Take for example the inner product of two length  $n$  vectors,  $x$  and  $y$ .

# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.
- Take for example the inner product of two length  $n$  vectors,  $x$  and  $y$ .

$$x^T y = \sum_i x_i y_i$$

# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.
- Take for example the inner product of two length  $n$  vectors,  $x$  and  $y$ .

$$x^T y = \sum_i x_i y_i$$

```
Input: x, y
Output: out
out = 0
for i in 1,...,n:
    out += x[i]*y[i]
```

# Inner Products

- A useful skill for evaluating numerical algorithms is translation between sums/products and loops.
- Take for example the inner product of two length  $n$  vectors,  $x$  and  $y$ .

$$x^T y = \sum_i x_i y_i$$

```
Input: x, y
Output: out
out = 0
for i in 1,...,n:
    out += x[i]*y[i]
```

A single sum/product can often be translated directly to a loop. This means that the inner product of two length  $n$  vectors has complexity  $\mathcal{O}(n)$

# Outer Products

Outer products, on the other hand, contain no sums. Instead, the output is an  $n \times n$  matrix.

# Outer Products

Outer products, on the other hand, contain no sums. Instead, the output is an  $n \times n$  matrix.

$$A = xy^T \text{ s.t. } A_{ij} = x_i y_j$$

# Outer Products

Outer products, on the other hand, contain no sums. Instead, the output is an  $n \times n$  matrix.

$$A = xy^T \text{ s.t. } A_{ij} = x_i y_j$$

```
Input: x, y
Output: A
out = 0
for i in 1,...,n:
    for j in 1,...,n:
        A[i,j] = x[i]*y[j]
```

# Outer Products

Outer products, on the other hand, contain no sums. Instead, the output is an  $n \times n$  matrix.

$$A = xy^T \text{ s.t. } A_{ij} = x_i y_j$$

```
Input: x, y
Output: A
out = 0
for i in 1,...,n:
    for j in 1,...,n:
        A[i,j] = x[i]*y[j]
```

Outer products have complexity  $\mathcal{O}(n^2)$ .



# Matrix Multiplication

Matrix multiplication can be decomposed into a bunch of vector products. Let  $A$  and  $B$  be  $n \times n$  matrices, then

# Matrix Multiplication

Matrix multiplication can be decomposed into a bunch of vector products. Let  $A$  and  $B$  be  $n \times n$  matrices, then

$$C = AB \text{ s.t. } C_{ij} = A_{i:}B_{:j} = \sum_k A_{ik}B_{kj}$$

# Matrix Multiplication

Matrix multiplication can be decomposed into a bunch of vector products. Let  $A$  and  $B$  be  $n \times n$  matrices, then

$$C = AB \text{ s.t. } C_{ij} = A_{i:}B_{:j} = \sum_k A_{ik}B_{kj}$$

```
Input: A, B
Output: C
for i in 1,...,n:
    for j in 1,...,n:
        C[i,j] = 0
        for k in 1,...,n:
            C[i,j] += A[i,k]*B[k,j]
```

# Matrix Multiplication

Matrix multiplication can be decomposed into a bunch of vector products. Let  $A$  and  $B$  be  $n \times n$  matrices, then

$$C = AB \text{ s.t. } C_{ij} = A_{i:}B_{:j} = \sum_k A_{ik}B_{kj}$$

```
Input: A, B
Output: C
for i in 1,...,n:
    for j in 1,...,n:
        C[i,j] = 0
        for k in 1,...,n:
            C[i,j] += A[i,k]*B[k,j]
```

The naive method for matrix multiplication has complexity  $\mathcal{O}(n^3)$ .

# Strassen's Algorithm

Modern linear algebra libraries actually use more advanced matrix multiplication algorithms based on **Strassen's Algorithm**.

# Strassen's Algorithm

Modern linear algebra libraries actually use more advanced matrix multiplication algorithms based on **Strassen's Algorithm**.

## Strassen's Algorithm: Basic Idea

- The naive matrix multiplication method requires 8 operations to multiply  $2 \times 2$  matrices, however, Strassen found a way to do it in 7 operations at higher cost per operation.


# Strassen's Algorithm

Modern linear algebra libraries actually use more advanced matrix multiplication algorithms based on **Strassen's Algorithm**.

## Strassen's Algorithm: Basic Idea

- The naive matrix multiplication method requires 8 operations to multiply  $2 \times 2$  matrices, however, Strassen found a way to do it in 7 operations at higher cost per operation.
- Strassen's algorithm recursively divides a matrix into  $2 \times 2$  matrices and applies Strassen's  $2 \times 2$  method.

# Strassen's Algorithm






# Strassen's Algorithm

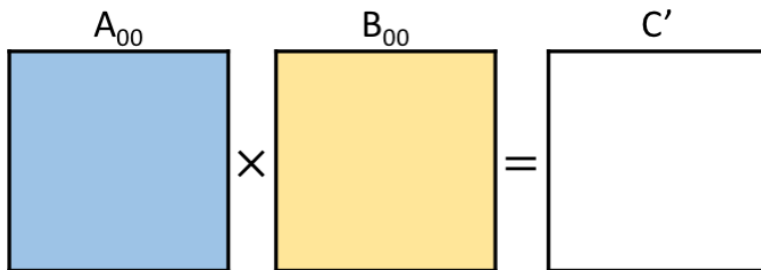
$$\begin{array}{|c|c|} \hline A & \\ \hline A_{00} & A_{01} \\ \hline A_{10} & A_{11} \\ \hline \end{array} \times \begin{array}{|c|c|} \hline B & \\ \hline B_{00} & B_{01} \\ \hline B_{10} & B_{11} \\ \hline \end{array} = \begin{array}{|c|c|} \hline C & \\ \hline C_{00} & C_{01} \\ \hline C_{10} & C_{11} \\ \hline \end{array}$$

# Strassen's Algorithm



$$C_{00} = A_{00}B_{00} + A_{01}B_{10}$$


# Strassen's Algorithm

$$A_{00} \times B_{00} = C'$$


# Strassen's Algorithm

$$\begin{array}{|c|c|} \hline A'_{00} & A'_{01} \\ \hline A'_{10} & A'_{11} \\ \hline \end{array} \times \begin{array}{|c|c|} \hline B'_{00} & B'_{01} \\ \hline B'_{10} & B'_{11} \\ \hline \end{array} = \begin{array}{|c|c|} \hline C'_{00} & C'_{01} \\ \hline C'_{10} & C'_{11} \\ \hline \end{array}$$

# Strassen's Algorithm



$$C'_{00} = A'_{00}B'_{00} + A'_{01}B'_{10}$$

# Strassen's Algorithm

- Strassen's algorithm reduces the complexity of matrix multiplication from  $\mathcal{O}(n^3)$  to  $\approx \mathcal{O}(n^{2.807})$ .

# Strassen's Algorithm

- Strassen's algorithm reduces the complexity of matrix multiplication from  $\mathcal{O}(n^3)$  to  $\approx \mathcal{O}(n^{2.807})$ .
- Many extensions to Strassen's have further reduced the complexity. Numpy uses one of these (see Demo).

# Strassen's Algorithm

- Strassen's algorithm reduces the complexity of matrix multiplication from  $\mathcal{O}(n^3)$  to  $\approx \mathcal{O}(n^{2.807})$ .
- Many extensions to Strassen's have further reduced the complexity. Numpy uses one of these (see Demo).
- Strassen's is typically faster for matrices of size  $100 \times 100$  and above. Good implementations will switch between naive and Strassen's.



# Demo

# Demo